

# epsml - ML Monorepo Proposal

Sep 29, 2025

Matt, Andrej, & Arjun

**TL;DR:** Our current codebase structure is causing compatibility issues across training and model serving. Let's migrate to a monorepo (with subpackages), and handle all dependencies using uv and python packaging.

[Current codebase structure](#)

[Proposed codebase structure](#)

[Appendix](#)

---

## Current codebase structure

[github.com/EpsilonLabsInc](https://github.com/EpsilonLabsInc)

- epsclassifiers
- epsutils
- epsdatasets
- 2d-image-encoders
- reports-pipeline
- ml-inference
- dinov3
- ...

### Problems:

1. **No python packaging** - In order to use code from one of the repos, you must clone the repo and add its local path to your `PYTHONPATH` environment variable. This is difficult to scale - spinning up a new VM or onboarding a new eng is painful.
2. **Outdated dependencies.** Not all repos have up-to-date `requirements.txt` files. It's not straightforward to simply create a virtual env and run `pip install -r requirements.txt`. Similarly, some `requirements.txt` have conflicts across repos, and we have no way of easily identifying/resolving this.
3. **Large builds due to broad imports** - Production code depends on `epsclassifiers`, which in turn depends on `epsutils`. `Epsutils` is a massive repo

with many dependencies. In order to run classifiers in production, we would need to include the entire (potentially conflict-prone) epsutils.

## Proposed codebase structure

[github.com/EpsilonLabsInc](https://github.com/EpsilonLabsInc)

```
→ epsml
    → pyproject.toml          # Workspace root \[details\]
    → uv.lock                 # Single lockfile for entire repo
    → .venv/                  # Single virtual env
    → packages/
        → epsclassifiers
            → pyproject.toml
            → src/epsclassifiers/
        → 2d-image-encoders
            → pyproject.toml
            → src/2d-image-encoders/
        → epsdatasets
            → pyproject.toml
            → src/epsdatasets/
        → epsanalytics
            → pyproject.toml
            → src/epsanalytics/
        → epstraining          # New, branched from epsutils
            → ...
        → epslabels           # ^ Same
        → epsdicom            # ^ Same
        → ...                  # Any future package goes here
→ epsutils                  # Deprecated, delete eventually
→ reports-pipeline
→ ml-inference                # References
→ dinov3
    → pyproject.toml          # External repos as packages.
    → uv.lock
→ InternVL-3x
    → pyproject.toml          # External repos as packages.
    → uv.lock
```

Internally, `uv` will ensure a single set of dependencies across all packages via a single lockfile, ensuring that we have no conflicts. Any changes within the monorepo will be reflected locally whenever developing internally - no rebuilding / repackaging would be necessary for other `epsml` packages [\[example\]](#).

External codebases (outside the monorepo) can add dependencies to the subpackages directly, reducing the need to add unnecessary dependencies. They will need to reference a monorepo package using an explicit tag (e.g. v0.1.0) or commit hash [\[example\]](#).

How to setup/use the monorepo:

```
Shell
# Clone monorepo
git clone https://github.com/EpsilonLabsInc/epsml
cd epsml

# Install uv
curl -LsSf https://astral.sh/uv/install.sh | sh

# Sync dependencies with lockfile (from git).
uv sync

# Run some code...
source .venv/bin/activate
cd packages/epsclassifiers
python src/epcclassifiers/run_training.py --dinov3 # example
```

### Advantages:

1. **Simple setup** - To set up a new machine / engineer, simply clone `epsml`, install `uv`, and run `uv sync`. Then, use the virtual environment to run your jobs.
2. **Easy dependencies for external code** - External repos (e.g. production inference code) will simply need to add a dependency of relevant packages within the monorepo, and use the monorepo lockfile to ensure a perfectly reproducible setup.
3. **Single lockfile guarantees no dependency conflicts** - `uv` will coordinate dependencies across all subpackages within the `epsml/packages/` directory. It will choose a single set of dependency versions to ensure no cross-package

compatibility issues.

4. **Consistent standards across all ML code** - We can define Github Actions for all of `epsilon_ml` that will run: dependency testing, pytesting, linting & formatting.

---

## Appendix

### A. Example monorepo pyproject.toml files

Example root pyproject.toml file:

```
None
# epsml/pyproject.toml

[tool.uv.workspace]
members = [
    "packages/*",
]
```

Example subpackage pyproject.toml files:

```
None
# epsml/packages/epsdatasets/pyproject.toml
dependencies = ["pandas>=2.0"]

# epsml/packages/epstraining/pyproject.toml
dependencies = ["torch>=2.0"]

# epsml/packages/epsclassifiers/pyproject.toml
dependencies = [
    "epsdatasets",
    "epstraining",
    "torch>=2.1"
]
```

## B. Example internal and external usage

Example 1: Testing changes within the monorepo (fixing a bug in epsdatasets and testing it in epsclassifiers)

Shell

```
# You're working in the epsml monorepo
cd ~/epsml/

# 1. Fix bug in epsdatasets
vim packages/epsdatasets/src/epsdatasets/loader.py
# Change line 42: return data.dropna() # Fixed the bug!

# 2. Update epsclassifiers to use the fix
vim packages/epsclassifiers/src/epsclassifiers/training.py
# Add: data = loader.load_data() # Uses the fixed version

# 3. Test immediately - no commits or tags needed
cd packages/epsclassifiers/
uv run python src/epsclassifiers/training.py
# ✅ Works! The bug fix is immediately visible

# 4. Run tests
uv run pytest
# ✅ All tests pass

# 5. Only NOW do you commit (once you're happy)
git add packages/epsdatasets packages/epsclassifiers
git commit -m "Fix data loading bug and update classifier"
git push

# 6. Tag a release when ready for external consumers
git tag v0.3.0
git push --tags
```

Example 2: Working outside the monorepo.

Shell

```
# You're working in ml-inference (external repo)
cd ~/ml-inference/
```

```
# Current state: ml-inference depends on epsclassifiers
cat pyproject.toml
```

```
None
# ml-inference/pyproject.toml
[project]
dependencies = ["epsclassifiers"]

[tool.uv.sources]
epsclassifiers = {
    git = "https://github.com/EpsilonLabsInc/epsml",
    subdirectory = "packages/epsclassifiers",
    tag = "v0.3.0" # ← Updated to the new version
}
```

## C. Eng notes

Arjun eng TODO:

- Git patch to move repo while maintaining the git history:  
<https://plankenau.com/blog/post/copy-commits-separate-git-repos>
- TODO: test dependencies work fine for the new setup (e.g. run inference on old setup vs new setup and ensure same outputs)
- Setup github actions for CI: run tests, run linter/formatter, run type checking, show results in PR

## Order of operations

- ~~1. [Matt/Andrej/Arjun] Confirm plan~~
- ~~2. [Arjun] Create epsml repo~~
- ~~3. [All] Ask everyone to submit existing code to their repos (main branch)~~

— ^ **Sept 30** ^ —

- ~~4. [Andrej] Follow [blog post](#) (to keep git history), copy in all source repos: epsclassifiers, epsdataset, epsutils, 2d-image-encoders, analytics~~

- a. ~~Then commit, without any restructuring / deleting~~
- 5. ~~[Andrej] Restructure — move epsutils subdirs into their own packages, delete old code, etc etc~~
  - a. ~~New structure: packages/epsclassifiers/src/epsclassifiers/...~~
  - b. ~~Include requirements.txt for easy uv dependencies later~~
  - c. ~~Cleanup~~
  - d. ~~Then commit~~

— ^ **Oct 1** ^ —

- 6. [Matt/Arjun] Address dependencies with uv
  - a. Look at requirements.txt, probably most are up-to-date but may be
- 7. [Arjun/Andrej] Test example classification training with DINOv3, ensure same/similar results as before

— ^ **Oct 2-5** ^ —

- 8. [All] DICOM code conflicts
- 9. [Arjun] Run code formatter over all files
  - a. [\[black python linter\]](#) or [\[ruff python linter\]](#)

— ^ **Oct 6-7** ^ —

- 10. Final step, add all remaining repos when we get time
  - a. 3d-image-encoders
  - b. ...
- 11. Set up Github Actions: dependency testing, pytesting, linting & formatting.